

Assignment 5

Classifying Movie Reviews using *Feature-based* Machine Learning & Deep Learning with RNN

1. Document the parameters you have used when defining all three models, and their respective obtained accuracies.

Part 1 – Using sci-kit klearn

My code for solving the problem can be seen [here](#), on Kaggle.

```
vector = HashingVectorizer(stop_words="english", binary=True)
```

I use stop_words = "English", to remove English stop words, and binary = True to set all nonzero counts to 1. "This is useful for discrete probabilistic models that model binary events rather than integer counts." – scikit-learn.org

```
'''DecisionTreeClassifier'''
dtt_classifier = DecisionTreeClassifier(criterion="entropy")
```

I used the parameter criterion = "entropy", because I wanted *information gain* instead of *Gini impurity*. The results are [similar](#), and I chose it since it's the one we have used in previous assignments. Entropy takes a little longer to compute, because of the logarithm.

$$Gini : Gini(E) = 1 - \sum_{j=1}^c p_j^2$$

$$Entropy : H(E) = - \sum_{j=1}^c p_j \log p_j$$

Figure 1: Equations for Entropy and Gini impurity

```
'''Bernulli Naive Bayes'''
bnb_classifier = BernoulliNB()
```

Nothing to comment here.

```
DecisionTreeClassifier    --> The accuracy for this classifier:
0.8645118288796274
BernoulliNB              --> The accuracy for this classifier:
0.8380500735474381
```

As we can see here, the Decision Tree Classifier and Bernoulli NB achieves **86.45%** and **83.80% accuracy**, respectively.

Part 2 – Using Keras

My code for solving the problem can be seen [here](#), on Kaggle.

Built the following neural network:

```
# initialize the Sequential model
model = Sequential()

# add layers to the model
model.add(Embedding(input_dim=vocab_size, output_dim=256,
input_length=max_length))
model.add(LSTM(256))
model.add(Dense(2, activation="sigmoid"))

# compile the model
# chose RMSProp as the optimizer, since it is usually a good choice for
recurrent neural networks.
model.compile(optimizer=RMSprop(), loss='binary_crossentropy',
metrics=['accuracy'])
```

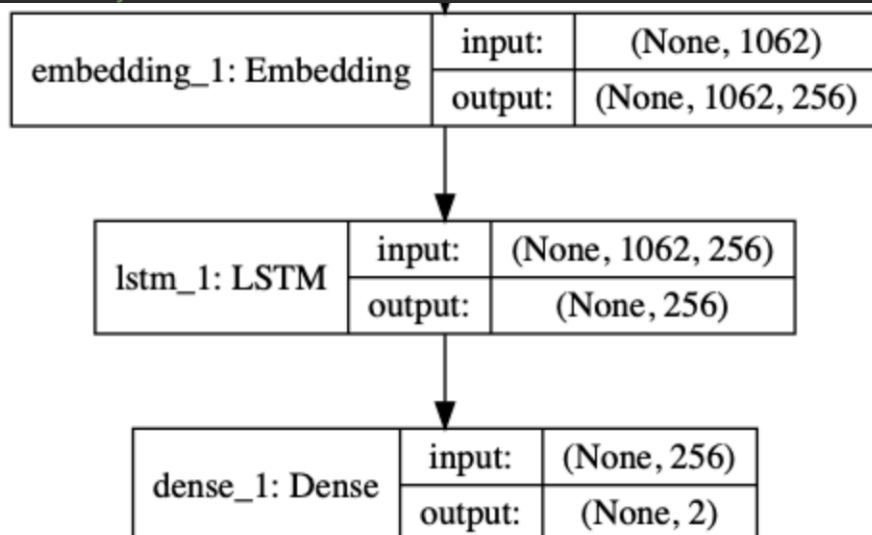


Figure 2: the LSTM model used in this assignment

For the dimension input- and output dimensions I had to do some “trial and error”. Used Stackoverflow and our friends over at “the Internet”, and watched how others adjusted their parameters to fit the training data.

A key factor was the amount of data we got to train on. It was **a lot** of data. So, I made the decision, after advice from the student assistant, to partition the data. I partitioned the data to fit the time gap I got at Kaggle (where I had a P100 to train on). Partitioned the data then by half the size of the X_train data.

In the code, I am using `to_categorical`, reason being that we have a Dense-layer outputting, through two “Densenodes”, a vector with two values, e.g.: $\rightarrow [[0.2, 0.8], [0.25, 0.75]]$ which is Densenode1’s and Densenode2’s probabilities for the the input being **class1** or **class2** (positive or negative «movie review»).

The RNN (with LSTM layer) achieved the following accuracy on the test data:

Loss: 0.1628359776079664

Accuracy: **0.9487619514586909**

Which I would say, is acceptable – considering we should acquire at least 90% accuracy.

Here is a graph for the training of the model used to achieve this result:




Figure 3: Evaluation of the model

2. What is the reason for the large improvement in accuracy from the Naive Bayes/Decision Tree models to the LSTM? Give the most important reasons

The biggest difference I believe, is that the LSTM network is capable of fetching the order of words, which makes a huge difference when the word “not” comes to play. “Nice” vs “not nice”. There’s a difference. The binarized vector of word-id’s that we send in to the embedded layer is also a far more efficient representation of a review than sparse matrices.